

Differentiable integer linear programming

ICLR'25

Zijie Geng , Jie Wang, Xijun Li, Fangzhou Zhu,
Jianye Hao, Bin Li, Feng Wu

ML for integer programming

Mixed integer linear programs (MILP):

- Flexible modeling tool for NP-hard combinatorial optimization
- E.g., scheduling, network design, ...
- Solvers are powerful but very computationally expensive

Challenge of ML-based heuristics (e.g., last class):

Supervision is expensive: requires solving NP-hard problems

This paper: unsupervised learning approach
via end-to-end differentiable pipeline

Overview of approach: DiffLO

$$\begin{aligned} & \text{minimize} && \mathbf{c}^\top \mathbf{x} \\ & \text{subject to} && A\mathbf{x} \leq \mathbf{b} \\ & && \mathbf{x} \in \{0, 1\}^n \end{aligned}$$

1. Relax to **probabilistic, continuous** equivalent form
2. Convert from **constrained** optimization to **unconstrained**
3. Reparameterize so objective is **differentiable** almost everywhere

1: Relax to probabilistic, continuous equivalent form

$$\begin{aligned} & \text{minimize } \mathbf{c}^T \mathbf{x} \\ & \text{subject to } A\mathbf{x} \leq \mathbf{b} \\ & \quad \mathbf{x} \in \{0, 1\}^n \end{aligned}$$



$$\begin{aligned} & \text{minimize } \mathbf{c}^T \hat{\mathbf{x}} \\ & \text{subject to } \hat{\mathbf{x}} \in [0, 1]^n \\ & \quad \mathbb{E}[\max\{A\hat{\mathbf{x}} - \mathbf{b}, 0\}] = \mathbf{0} \end{aligned} \quad \xrightarrow{\quad x_i \sim \text{Bernoulli}(\hat{x}_i) \quad}$$

Justification of probabilistic form:

- **Thm 1** (informal): top is feasible & solvable iff bottom is too
- **Thm 2** (informal): opt solution of top \equiv (rounded) solutions of bottom

Overview of approach: DiffLO

$$\begin{aligned} & \text{minimize} && \mathbf{c}^\top \mathbf{x} \\ & \text{subject to} && A\mathbf{x} \leq \mathbf{b} \\ & && \mathbf{x} \in \{0, 1\}^n \end{aligned}$$

1. Relax to **probabilistic, continuous** equivalent form
2. Convert from **constrained** optimization to **unconstrained**
3. Reparameterize so objective is **differentiable** almost everywhere

2: Convert from constrained to unconstrained

$$\begin{aligned} & \text{minimize } \mathbf{c}^\top \hat{\mathbf{x}} \\ & \text{subject to } \hat{\mathbf{x}} \in [0, 1]^n \\ & \quad \mathbb{E}[\max\{A\mathbf{x} - \mathbf{b}, \mathbf{0}\}] = \mathbf{0} \end{aligned}$$

\mathbf{a}_j : jth row of A

$\hat{\phi}_j(\hat{\mathbf{x}}) = \mathbb{E}_{\mathbf{x} \sim p(\cdot | \hat{\mathbf{x}})} [\max\{\mathbf{a}_j^\top \mathbf{x} - b_j, 0\}]$: expected violation of jth constraint
Independent Bernoullis

$$\begin{aligned} & \text{minimize } \mathbf{c}^\top \hat{\mathbf{x}} + \mu \sum_{j=1}^m \hat{\phi}_j(\hat{\mathbf{x}}) \\ & \text{subject to } \hat{\mathbf{x}} \in [0, 1]^n \end{aligned}$$

Overview of approach: DiffILO

$$\begin{aligned} & \text{minimize} && \mathbf{c}^\top \mathbf{x} \\ & \text{subject to} && A\mathbf{x} \leq \mathbf{b} \\ & && \mathbf{x} \in \{0, 1\}^n \end{aligned}$$

1. Relax to **probabilistic, continuous** equivalent form
2. Convert from **constrained** optimization to **unconstrained**
3. Reparameterize so objective is **differentiable** almost everywhere

3: Reparameterize so objective is differentiable a.e.

Challenge to applying SGD: $\nabla_{\hat{\mathbf{x}}}\hat{\phi}_j(\hat{\mathbf{x}}) = \nabla_{\hat{\mathbf{x}}}\mathbb{E}_{\mathbf{x} \sim p(\cdot|\hat{\mathbf{x}})}[\max\{\mathbf{a}_j^\top \mathbf{x} - b_j, 0\}]$

3: Reparameterize so objective is differentiable a.e.


Challenge to applying SGD: $\nabla_{\hat{\mathbf{x}}}\hat{\phi}_j(\hat{\mathbf{x}}) = \nabla_{\hat{\mathbf{x}}}\mathbb{E}_{\mathbf{x} \sim p(\cdot|\hat{\mathbf{x}})}[\max\{\mathbf{a}_j^\top \mathbf{x} - b_j, 0\}]$

Very messy to differentiate!

3: Reparameterize so objective is differentiable a.e.

Challenge to applying SGD: $\nabla_{\hat{\mathbf{x}}}\hat{\phi}_j(\hat{\mathbf{x}}) = \nabla_{\hat{\mathbf{x}}}\mathbb{E}_{\mathbf{x} \sim p(\cdot|\hat{\mathbf{x}})}[\max\{\mathbf{a}_j^\top \mathbf{x} - b_j, 0\}]$
Very messy to differentiate!

Instead: relax to continuous RV $\xi(\hat{x}_i)$ such that $\mathbb{P}[\xi(\hat{x}_i) > 0.5] = \hat{x}_i$




3: Reparameterize so objective is differentiable a.e.

Challenge to applying SGD: $\nabla_{\hat{\mathbf{x}}}\hat{\phi}_j(\hat{\mathbf{x}}) = \nabla_{\hat{\mathbf{x}}}\mathbb{E}_{\mathbf{x} \sim p(\cdot|\hat{\mathbf{x}})}[\max\{\mathbf{a}_j^\top \mathbf{x} - b_j, 0\}]$

Very messy to differentiate!

Instead: relax to continuous RV $\xi(\hat{x}_i)$ such that $\mathbb{P}[\xi(\hat{x}_i) > 0.5] = \hat{x}_i$

1. Apply logit function $\tau(\hat{x}_i) = \log \frac{\hat{x}_i}{1 - \hat{x}_i}$ (inverse of sigmoid)



3: Reparameterize so objective is differentiable a.e.

Challenge to applying SGD: $\nabla_{\hat{\mathbf{x}}}\hat{\phi}_j(\hat{\mathbf{x}}) = \nabla_{\hat{\mathbf{x}}}\mathbb{E}_{\mathbf{x} \sim p(\cdot|\hat{\mathbf{x}})}[\max\{\mathbf{a}_j^\top \mathbf{x} - b_j, 0\}]$
Very messy to differentiate!

Instead: relax to continuous RV $\xi(\hat{x}_i)$ such that $\mathbb{P}[\xi(\hat{x}_i) > 0.5] = \hat{x}_i$

1. Apply logit function $\tau(\hat{x}_i) = \log \frac{\hat{x}_i}{1 - \hat{x}_i}$ (inverse of sigmoid)
2. Sample $\epsilon \sim U(0, 1)$



3: Reparameterize so objective is differentiable a.e.

Challenge to applying SGD: $\nabla_{\hat{\mathbf{x}}}\hat{\phi}_j(\hat{\mathbf{x}}) = \nabla_{\hat{\mathbf{x}}}\mathbb{E}_{\mathbf{x} \sim p(\cdot|\hat{\mathbf{x}})}[\max\{\mathbf{a}_j^\top \mathbf{x} - b_j, 0\}]$

Very messy to differentiate!

Instead: relax to continuous RV $\xi(\hat{x}_i)$ such that $\mathbb{P}[\xi(\hat{x}_i) > 0.5] = \hat{x}_i$

1. Apply logit function $\tau(\hat{x}_i) = \log \frac{\hat{x}_i}{1 - \hat{x}_i}$ (inverse of sigmoid)
2. Sample $\epsilon \sim U(0, 1)$
3. Perturb logit: $\tau(\hat{x}_i) + \tau(\epsilon)$




3: Reparameterize so objective is differentiable a.e.

Challenge to applying SGD: $\nabla_{\hat{\mathbf{x}}}\hat{\phi}_j(\hat{\mathbf{x}}) = \nabla_{\hat{\mathbf{x}}}\mathbb{E}_{\mathbf{x} \sim p(\cdot|\hat{\mathbf{x}})}[\max\{\mathbf{a}_j^\top \mathbf{x} - b_j, 0\}]$

Very messy to differentiate!

Instead: relax to continuous RV $\xi(\hat{x}_i)$ such that $\mathbb{P}[\xi(\hat{x}_i) > 0.5] = \hat{x}_i$

1. Apply logit function $\tau(\hat{x}_i) = \log \frac{\hat{x}_i}{1 - \hat{x}_i}$ (inverse of sigmoid)
2. Sample $\epsilon \sim U(0, 1)$
3. Perturb logit: $\tau(\hat{x}_i) + \tau(\epsilon)$
4. Map back to (0,1): $\xi(\hat{x}_i; \epsilon) = \sigma(\tau(\hat{x}_i) + \tau(\epsilon))$




3: Reparameterize so objective is differentiable a.e.

Challenge to applying SGD: $\nabla_{\hat{\mathbf{x}}}\hat{\phi}_j(\hat{\mathbf{x}}) = \nabla_{\hat{\mathbf{x}}}\mathbb{E}_{\mathbf{x} \sim p(\cdot|\hat{\mathbf{x}})}[\max\{\mathbf{a}_j^\top \mathbf{x} - b_j, 0\}]$

Very messy to differentiate!

Surrogate that's differentiable almost everywhere:

$$\mathbb{E}_{\mathbf{x} \sim p(\cdot|\hat{\mathbf{x}})}[\max\{\mathbf{a}_j^\top \mathbf{x} - b_j, 0\}] \approx \mathbb{E}_\epsilon[\max\{\mathbf{a}_j^\top \xi(\hat{\mathbf{x}}; \epsilon) - b_j, 0\}] := \hat{\varphi}_j(\hat{\mathbf{x}})$$



Graph neural network

$$\begin{aligned} & \text{minimize} && \mathbf{c}^\top \mathbf{x} \\ & \text{subject to} && A\mathbf{x} \leq \mathbf{b} \\ & && \mathbf{x} \in \{0, 1\}^n \end{aligned}$$

Represent IP with a constraint-variable bipartite graph \mathcal{G} (like last class)

Graph neural network

$$\begin{aligned} & \text{minimize} && \mathbf{c}^\top \mathbf{x} \\ & \text{subject to} && A\mathbf{x} \leq \mathbf{b} \\ & && \mathbf{x} \in \{0, 1\}^n \end{aligned}$$

Represent IP with a constraint-variable bipartite graph \mathcal{G} (like last class)

$$\text{GNN } f_\theta(\mathcal{G}) = \hat{\mathbf{x}} \in [0, 1]^n$$

Graph neural network

$$\begin{aligned} & \text{minimize} && \mathbf{c}^\top \mathbf{x} \\ & \text{subject to} && A\mathbf{x} \leq \mathbf{b} \\ & && \mathbf{x} \in \{0, 1\}^n \end{aligned}$$

Represent IP with a constraint-variable bipartite graph \mathcal{G} (like last class)

$$\text{GNN } f_\theta(\mathcal{G}) = \hat{\mathbf{x}} \in [0, 1]^n$$

$$\text{Loss function } \mathcal{L}(\theta; \mathcal{G}) = \mathbf{c}^\top f_\theta(\mathcal{G}) + \mu \sum_{j=1}^m \hat{\varphi}_j(f_\theta(\mathcal{G}))$$

Inference


Sample from Bernoullis $\mathbf{x}' \sim p(\cdot \mid f_\theta(\mathcal{G}))$

Solve (e.g., with Gurobi):

$$\begin{aligned} & \text{minimize } \mathbf{c}^\top \mathbf{x} \\ & \text{subject to } A\mathbf{x} \leq \mathbf{b} \\ & \quad \sum_{i:x'_i=0} x_i + \sum_{i:x'_i=1} (1 - x_i) \leq \Delta \\ & \quad \mathbf{x} \in \{0, 1\}^n \end{aligned}$$

Training time comparison

240 IPs for training, 60 for validation, 100 for testing



Objective values

	SC (min, BKS: 86.45)			IS (max, BKS:684.14)			CA (max, BKS:22272.55)		
	10s	100s	1000s	10s	100s	1000s	10s	100s	1000s
Gurobi	1031.39	<u>87.09</u>	<u>86.52</u>	682.02	<u>684.12</u>	<u>684.13</u>	22090.76	22242.58	22272.03
PS+Gurobi	<u>131.87</u>	125.26	125.26	684.13	684.13	<u>684.13</u>	22140.65	<u>22243.12</u>	<u>22272.47</u>
DiffILO+Gurobi	95.65	86.78	86.48	<u>684.00</u>	684.12	684.14	22177.82	22260.48	22272.55

Overview

- **Goal:** Learn to solve IPs without supervision or solver labels
 - a. Reformulate discrete IP as continuous, probabilistic program
 - b. Add exact penalty to remove constraints
 - c. Apply relaxed Bernoulli for differentiable sampling
- Resulting objective differentiable almost everywhere
- Unsupervised: fast training
- Improves solver warm starts